




Assignment 1: Sentiment Classification of Twitter Tweets (NLP Pipeline)

 **Difficulty:** Advanced |  **Time:** 4–5 hours

 **Dataset:** Twitter US Airline Sentiment

 **Tools:** Python, NLTK/spaCy, Scikit-learn, Pandas, NumPy, Seaborn, Matplotlib, Streamlit or MLflow

Problem Statement

You are working for a customer support analytics team at an airline company. Your task is to **build a sentiment classifier** that can automatically categorize tweets about airline services into **positive, negative, or neutral**. The solution should be modular and support continuous improvements and deployment.

Assignment Objectives

1. Build a complete ML pipeline to classify tweets using TF-IDF + traditional ML models (Logistic Regression / Naive Bayes).
2. Use OOP principles to structure code into reusable classes/functions.
3. Deploy the model with a frontend using Streamlit or register and track it using MLflow.

Task Breakdown

Task 1: Data Cleaning & Exploration

- Load and explore the dataset
- Handle missing values
- Visualize class imbalance

Task 2: Preprocessing with OOP

- Create a class TweetPreprocessor:
 - `clean_text()`: remove mentions, links, emojis
 - `tokenize_and_lemmatize()`: using spaCy or NLTK
 - `remove_stopwords()` method

Task 3: Modeling Pipeline

- Create a class SentimentModel:
 - `vectorize()` using TF-IDF
 - `train_model()` using Logistic Regression
 - `evaluate()` using F1, accuracy, confusion matrix

Task 4: Deployment

Option A – Streamlit

- User enters tweet → predicted sentiment is shown



Option B – MLflow


- Track model:
 - Preprocessing steps
 - Model accuracy
 - Parameters
- Register best model


Deliverables

- Python notebook or scripts using modular OOP
- TF-IDF vectorizer saved
- Deployment (Streamlit or MLflow)
- Screenshot of model evaluation or UI
- README with setup and usage

Assignment 2: News Topic Classification Using BERT (Transformer-based NLP)

 **Difficulty:** Expert |  **Time:** 5–6 hours

 **Dataset:** AG News Classification Dataset on Kaggle

 **Tools:** HuggingFace Transformers, PyTorch or TensorFlow, Pandas, Seaborn, Matplotlib, Streamlit or MLflow

Problem Statement

A media analytics company wants to categorize news articles automatically into one of 4 categories: **World, Sports, Business, and Sci/Tech**. You are asked to build a transformer-based pipeline using **BERT** to improve classification accuracy over traditional methods.

Assignment Objectives

1. Use HuggingFace's BERT model for fine-tuning on multi-class classification.
2. Implement reusable classes/functions for tokenization, modeling, and prediction.
3. Integrate experiment tracking or UI deployment using MLflow or Streamlit.

Task Breakdown

Task 1: Data Loading and Cleaning

- Load train/test CSVs
- Basic text cleanup (remove special characters, optional)

Task 2: Tokenization & Encoding

- Create class NewsTokenizer:
 - Load BERT tokenizer

- Tokenize and pad sequences
- Encode target labels

Task 3: Model Training

- Create class `NewsClassifier`:
 - Load `bert-base-uncased`
 - Freeze/unfreeze layers
 - Train with learning rate scheduling and validation loop

Task 4: Evaluation & Metrics

- Use F1 score, confusion matrix, and classification report

Task 5: Deployment

Option A – Streamlit

- Upload a text → category is predicted
- Dropdown to select pre-trained or fine-tuned model

Option B – MLflow


- Log:
 - Training time
 - Accuracy, F1
 - Model name, tokenizer, parameters
- Save and register model artifact

Deliverables


- Scripts: `tokenizer.py`, `bert_model.py`, `utils.py`
- Saved model and tokenizer
- Streamlit or MLflow deployment
- Instructions in README

Bonus: Instructions Format for Submission


Each assignment should follow this directory structure:

 assignment_name/


|


├──  data/ # Store dataset here

├──  models/ # Saved models or tokenizer

├──  app/ # Streamlit or Flask app (optional)

├──  main.py # Main runner file

├──  README.md # Setup, usage, deployment instructions

└──  requirements.txt # List of dependencies